

Data ~~Management~~ Compression for Data Lakes

Jonathan WINANDY
primatice.com

About ME

- I am (Not Only) a Data Engineer

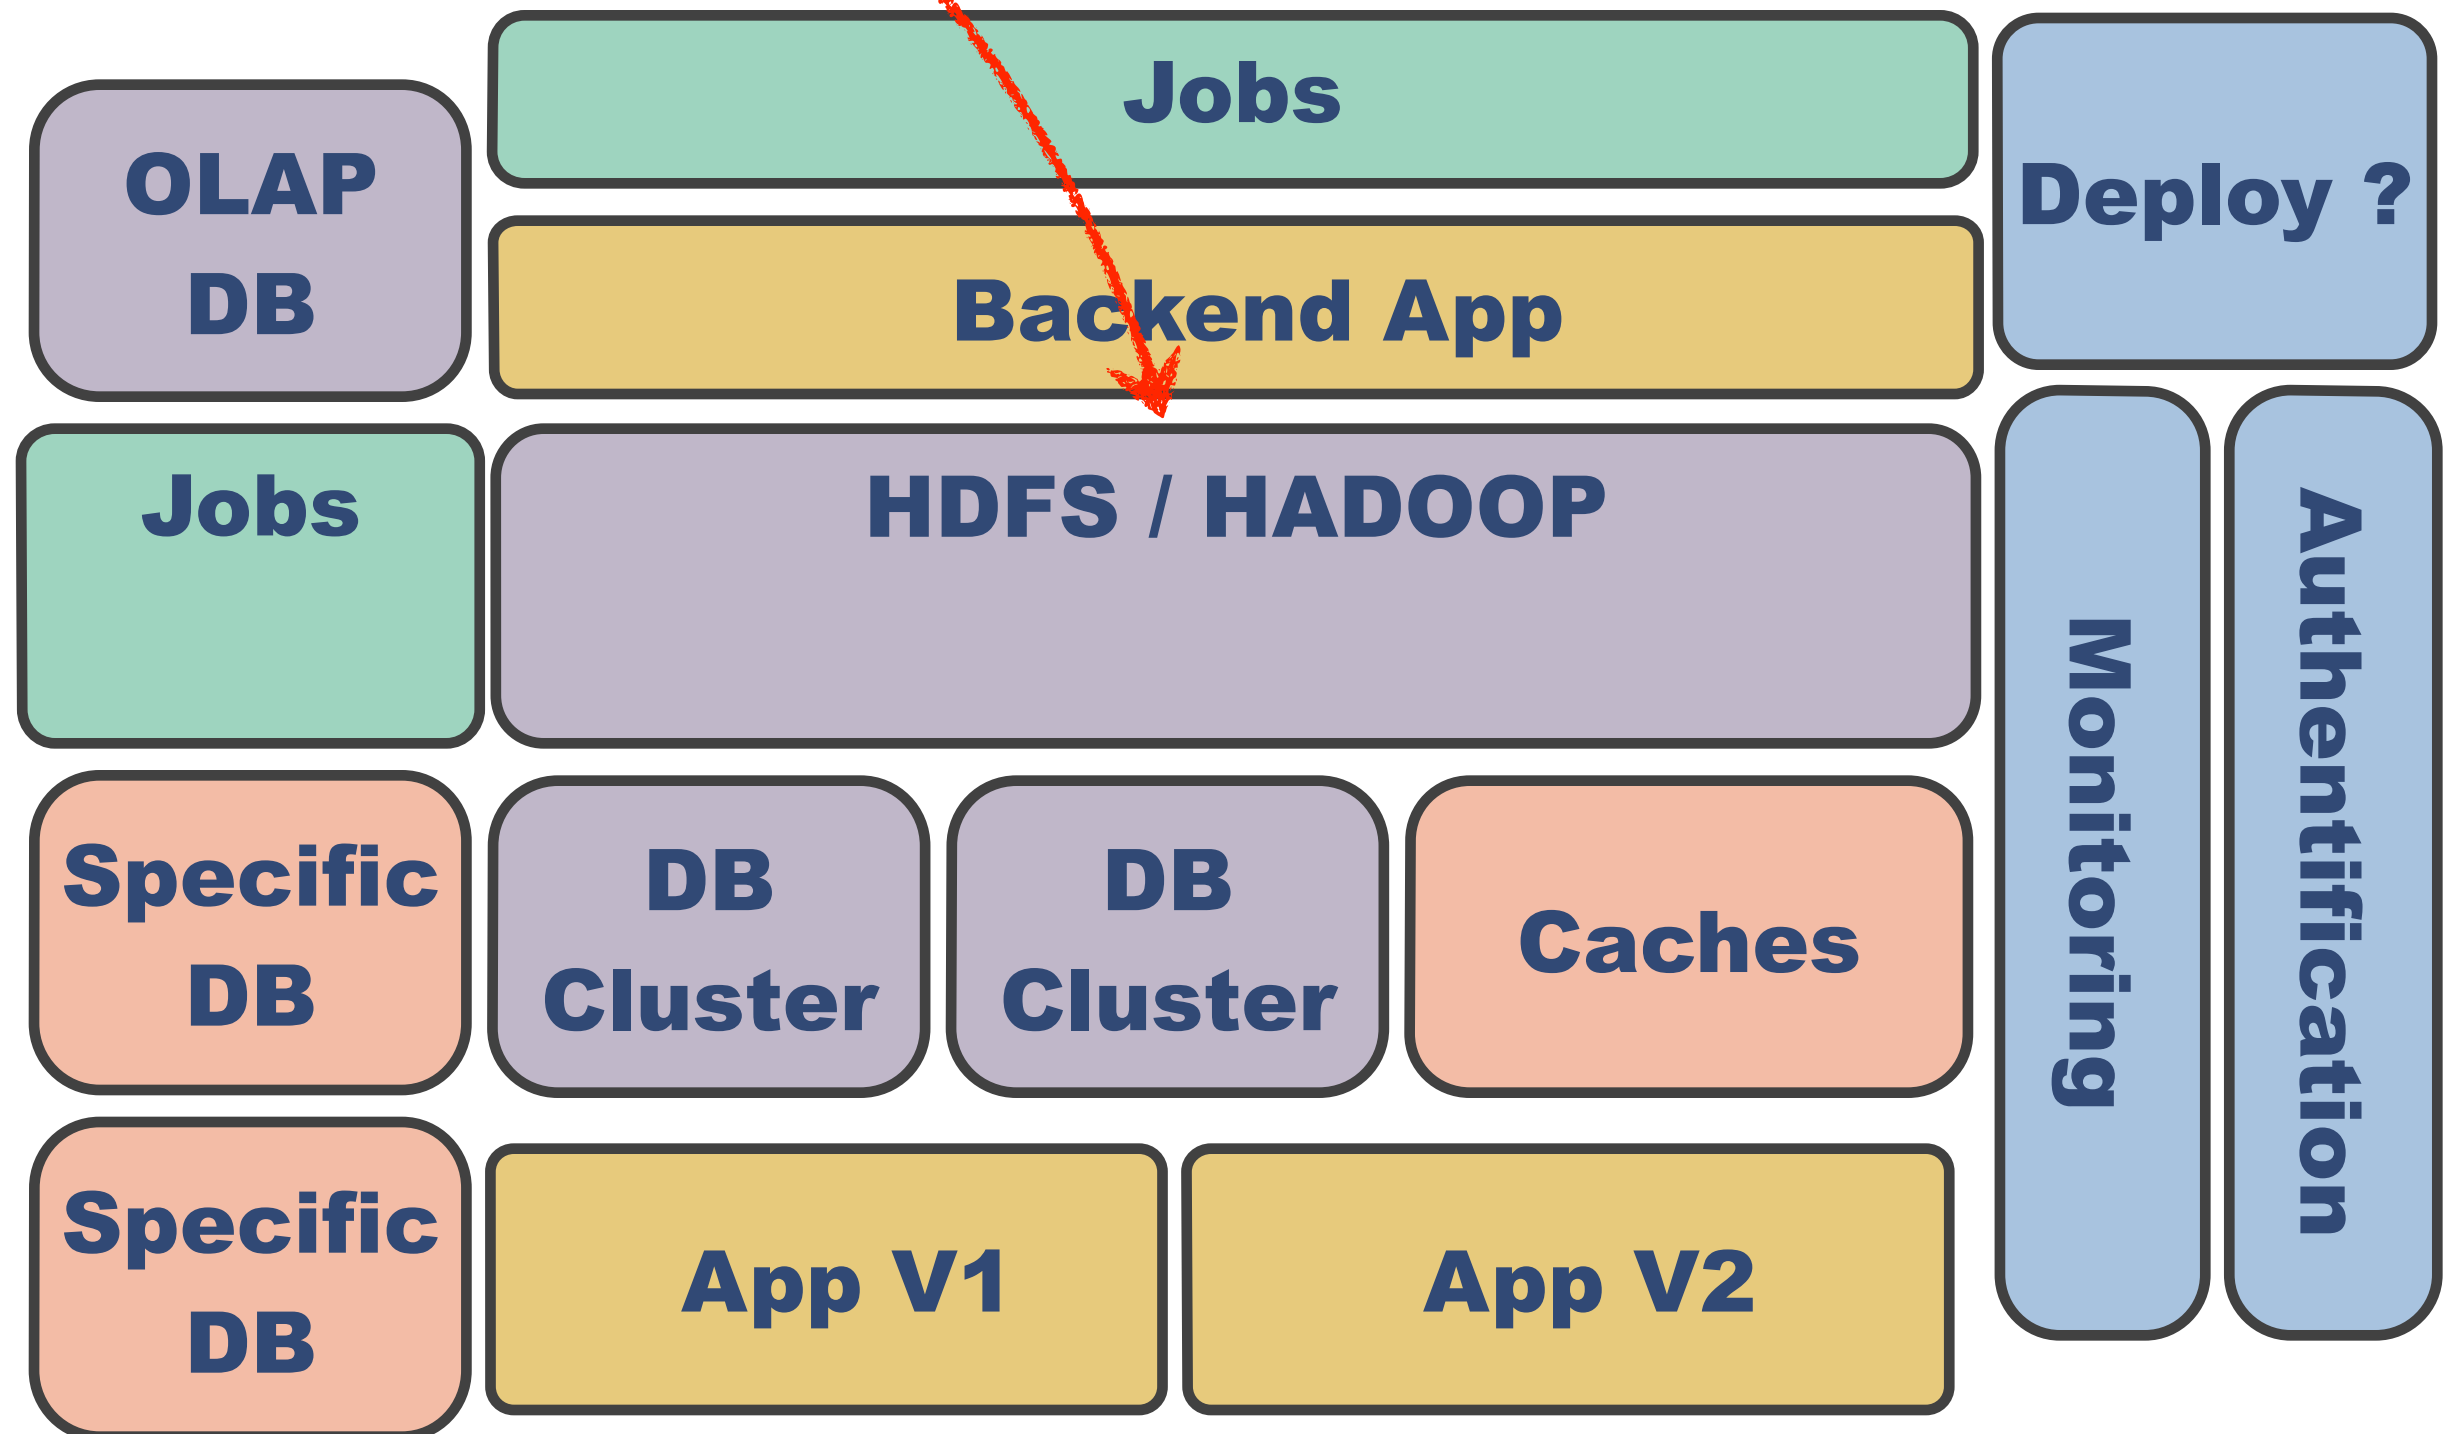
Contents

- Data Lake : What is a Data Lake ?
- Compute : What can we do with the gathered data ?
- Conclusion !

Data Lake

Infrastructure > The Data Lake*

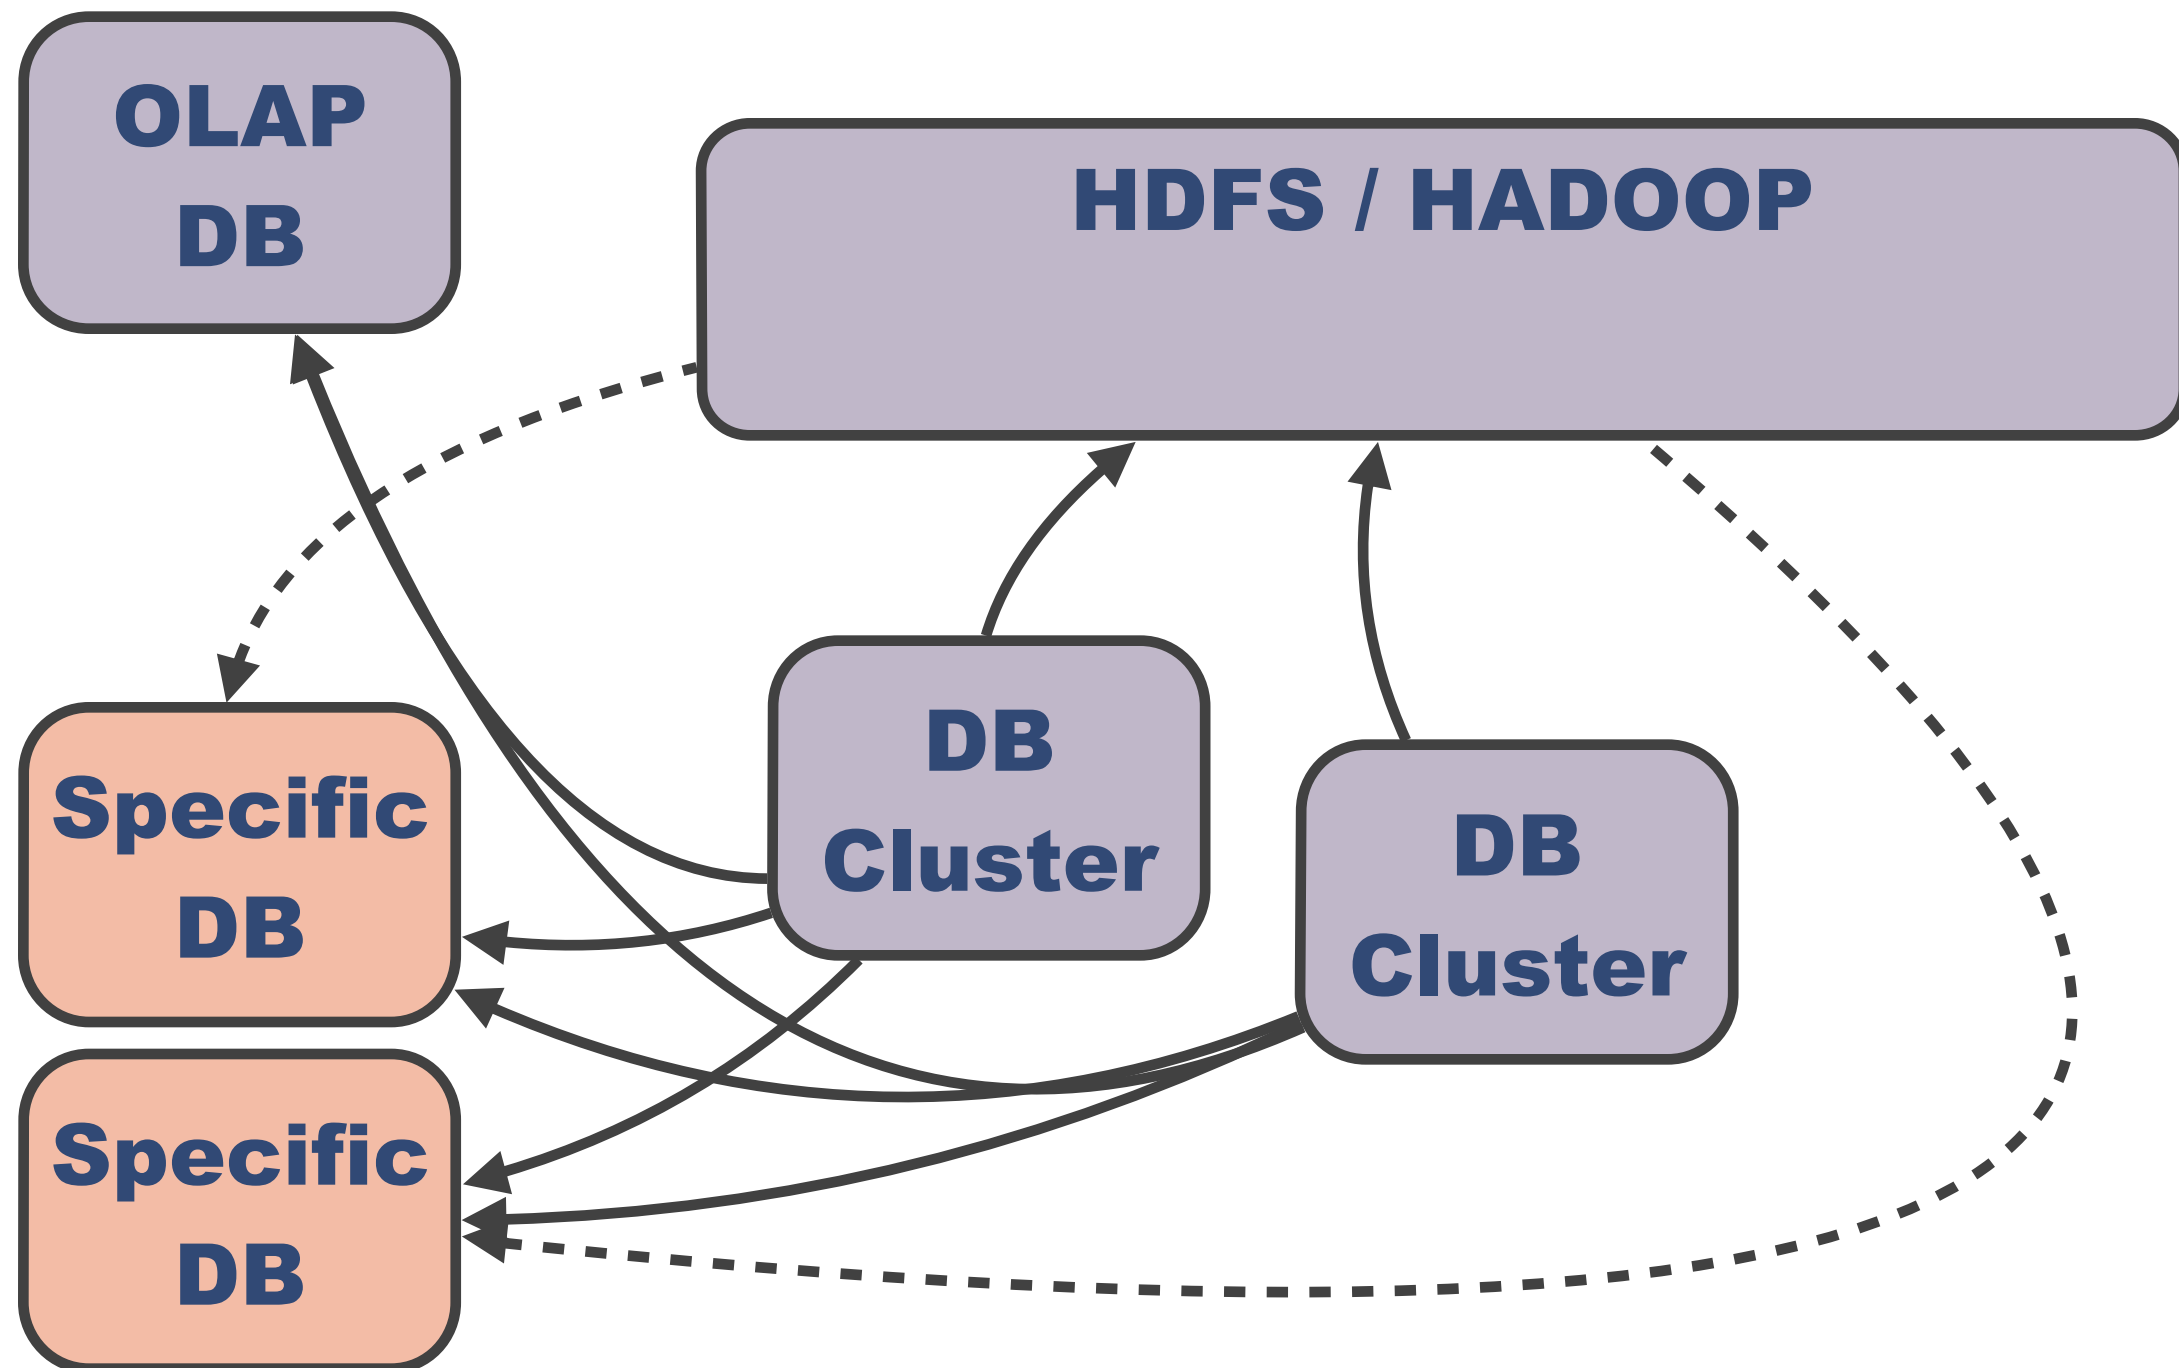
(*) you actually need more than just Hadoop to make a data lake.



Data Lake


We use HDFS to store a copy of all Data.

And we use HADOOP to create views on this Data.




Offloading to HDFS

Offloading data is just making a copy of databases data onto HDFS as a directory with a couple of files (parts).



Data Lake

For relational databases, we can use apache sqoop to copy for example the table ***payments*** from schema ***sales*** into ***hdfs://staging/sales/payments/2014-09-22T18_01_10***
into ***hdfs://staging/sales/payments/2014-09-21T17_50_25***
into ***hdfs://staging/sales/payments/2014-09-20T18_32_43***



A light capacity planning for your potential cluster :

- If you have 100 GB of binary compressed data*,
- With a cost of storage around 0,03\$/per GB/per month,
- An offload every day would cost 90\$/per month of kept Data,
- In six month, this offloading would cost 1900\$ and weight 18TB (<- this is ***Big Data***).

* obviously not plain text/JSON

But for this 18TB, you have quite a lot of features :

- You can track and understand bugs and data corruption,
- Analysing the business without harming your production DBs,
- Bootstrap new products based on your existing Data,
- And also have now a real excuse to learn Hadoop or make some use of your existing Hadoop bare metal clusters !

Compute

Having a couple of snapshots in HDFS,
we can use the tremendous power of MapReduce
to join over the snapshots, and compute what
changed in the database during a day.

$$\text{delta} = f(\text{J}, \text{J} - 1)$$

$$\text{merge} = f(\text{J}, \text{J} - 1, \text{J} - 2, \text{J} - \dots, \dots)$$

Compute

Δ function takes 2 (or +) snapshots and output a merge view of those snapshots with an extra column 'dm'.

Inputs :

id	name
1	Jon
2	Bob
3	James

id	name
1	John
3	James
4	Jim

Output :

id	name	dm
1	John	01
1	Jon	10
3	James	11
4	Jim	01
2	Bob	10

<https://github.com/viadeo/viadeo-avro-utils>

Compute

Δ function takes 2 (or +) snapshots and output a merge view of those snapshots with an extra column 'dm'.

Inputs (generalised) :

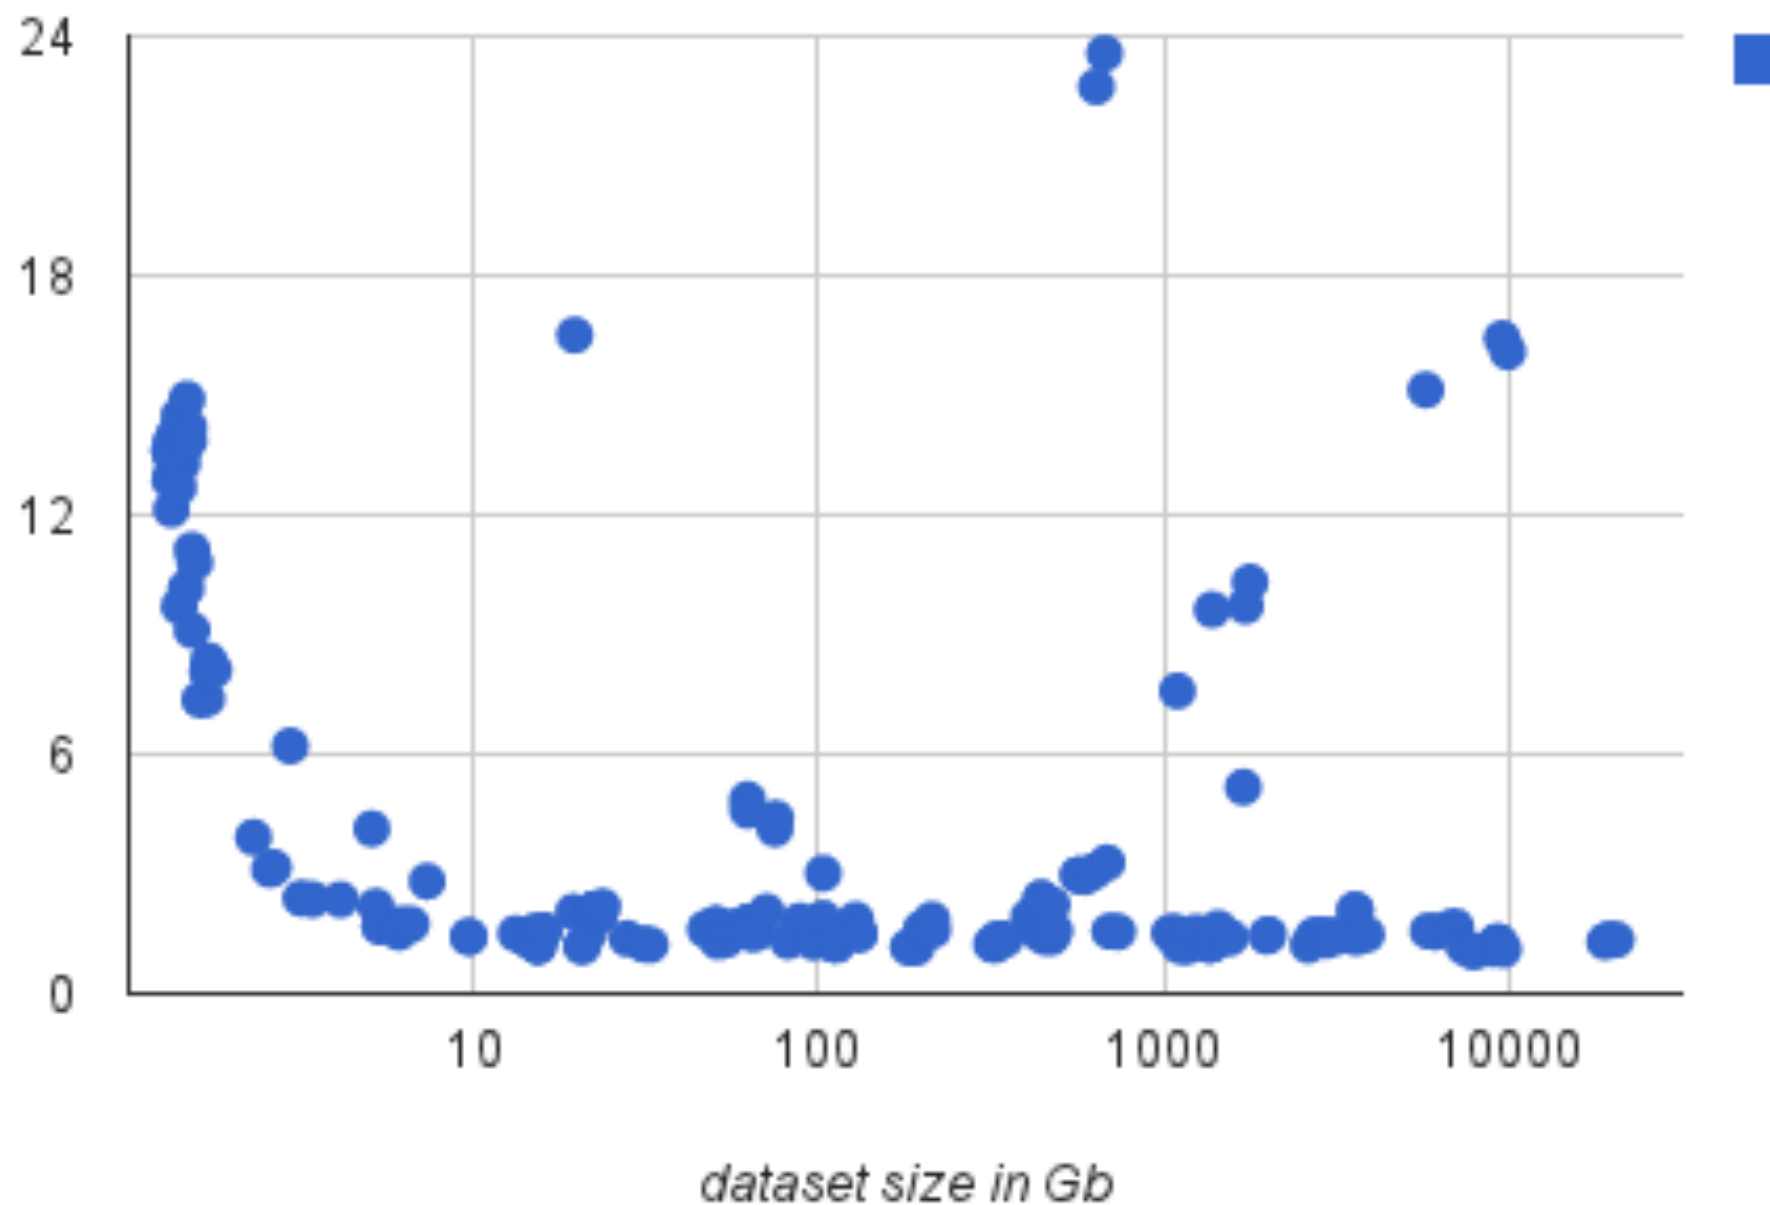
Output :

Δ (id	,	id	,	id	,	id)	=	1	1111
	1		1		1		1			2	1000
	2		3		4		3			3	1101
	3		4		5		5			4	0110
										5	0011

<https://github.com/viadeo/viadeo-avro-utils>

Analysing the compression : merging 50 snapshots

Ratio : Size of merge output / Size of largest input
Less is better, ≤ 1 is perfect




Compute

Analysing the compression : merging 50 snapshots

Ratio : Size of merge output / Size of largest input
Less is better, ≤ 1 is perfect

Small datasets
get a bit of overhead



OK, Just x2

WINS !!

Conclusion

- So now the offloading might cost just 36\$ of storage for 6 month (6\$ per month).
- Welcome back at small data !
- But what does this mean for production Data ?
Did you really need to store that in a mutable database in the first place ?

Thank you,
Questions ?